


# Dataflow Architectures

Karin Strauss

# Introduction

- Dataflow machines: *programmable* computers with hardware optimized for *fine grain data-driven* parallel computation
- **fine grain**: at the instruction granularity
- **data-driven**: activated by data availability
  - only data dependences constrain parallelism
  - programs usually represented as graphs


# Terminology



- nodes (FUs)
- tokens (data)
- arcs (dependences)
- input port
- enabling
- firing
- input data + firing = output data


# Node Types

- functional: (e.g. +, -, \*....)



- conditional:


- merge:



non-strict firing rule  
acts as a serializer

# if ... then ... else ...

```
if(u < v)
    z = x + y;
else
    z = x - y;
```




# Iteration, Recursion, Reuse and Reentrance

- using the *same graph* to perform computation on *different data sets*
- assume no storage elements are used, data is only present in wires

# while example

suppose this node takes 10x what a regular node takes to perform computation



t	x	y
0	m	m
1	<,c	c
2	c,c	c
3	+1	+
4	h,m	+
5	<,c	+
6	c,c	+
7	+1	+
8	h,m	+

# Synchronization

- static
  - locks (compound branch and merge nodes)
    - nodes only fire when all inputs are ready
    - loss of concurrency
  - acknowledging (control flow protocol)
    - extra arcs from consumer to producer
    - increases resources needed
      - can be reduced with detailed analysis

# Synchronization II

- dynamic
  - each iteration is executed in a separate instance of the graph
  - code copying
    - new instance of subgraph is created per iteration
    - need to direct tokens from earlier iterations to inputs of new iteration
  - tagged tokens
    - attach a tag to each token, associating it with an iteration
    - fire when input tokens have all the same tag

# Tagged Tokens


- create other problems
  - how to manage tags (size, distribution)
  - storage overhead
    - tags have to be stored with tokens
    - tokens that cannot be consumed at the moment may need to be stored for later use
  - too much parallelism!
    - storage overflow
    - deadlocks

# Procedures


- procedures can be called from several distinct calling sites
  - callee node address sent in special token
  - mechanism to direct the “return value” token

# Processing Element Architecture

- several processing elements (PEs) that communicate with each other




# Tagged Architectures



- 1) matching unit receives token
- 2) check memory; if all other inputs with same tag are there, send all tokens to fetching unit
- 3) fetching unit retrieves node from memory
- 4) fetch unit assembles an executable packet and sends it to functional unit
- 5) functional unit executes node with inputs provided by packet
- 6) output is sent back to matching unit


# One-level Architecture

- a functional unit delivers tokens to the enabling unit of the correct processing element




# Two-level Architecture

- each functional unit consists of several functional elements that can process packets in parallel




# Two-stage Architecture

- each enabling unit can send executable packets to any functional unit
  - good for heterogeneous functional units



# Manchester Dataflow Machine

- Gurd and Watson (1976-1981), two-level machine



- pipelined ring
- matching unit pairs tokens
- large data sets overflow to overflow units
- appropriate instruction is fetched from instruction store
- inputs and instruction are forwarded to processing unit

# Underutilization

- poor performance
- underutilization of functional units
  - imbalance
  - overhead computation
- underutilization of storage space
  - large data sets
  - code replication
  - tags, destination addresses

# Dataflow Model

- Benefits:
  - exposes parallelism
  - can tolerate latencies
  - mechanisms for fine-grain synchronization
- Drawbacks:
  - loss of locality (interleaving of instructions)
  - waste of resources
  - space overhead


# Trends

- convergence of dataflow architectures with conventional Von-Neumann architectures
- Ianucci's hybrid approach
- decoupled architectures
- out-of-order processors

# Ianucci's Hybrid Approach

- scheduling quanta
  - little or no parallelism among instructions
  - maximize run length (more locality)
  - minimize # of arcs between quanta
  - increase resource utilization
  - deadlock avoidance: dependence sets

# Ianucci's Hybrid Approach (II)



# Decoupled Architectures and Out-of-Order Processors

- Decoupled architectures:
  - decentralized structures (distinct FUs)
  - instruction steering based on input and output dependences, and operation to be performed
- Out-of-Order processors:
  - register renaming to identify “iteration”
  - instruction scheduling based on ready input operands
  - predication?

# Less Traditional Proposals

- Wavescalar (U. Washington, Mark Oskin)
  - PIM architecture
  - supports traditional Von-Neumann style memory semantics in a dataflow model
  - any programming language
- Edge/Trips (U.T. Austin, Doug Burger)
  - direct instruction communication (within blocks)
  - groups of 128 instructions mapped to an array of execution units: dataflow within, sequential across
  - loads and stores still do through memory ordering hw

# Related OoO Techniques

- instruction collapsing (Micro'37)
  - strands: dependent instructions with intermediate computation that does not need to be committed to architectural state (Wills - Georgia Tech)
    - e.g.  $e = a + (b + (c + d))$
  - mini-graphs: sequence of instructions with at most 2 inputs, 1 output, one memory operation and 1 terminal control transfer (Roth - U. Pennsylvania)
  - goal: save processor resources
    - instruction queue entries
    - reorder buffer entries
    - registers / register file accesses